

WHEN DOES TRAINING ON DOWNSCALED IMAGES YIELD THE SAME GRADIENTS?

Seunghyun Ji

Independent researcher

standingbehindnv@gmail.com

ABSTRACT

Scale-wise diffusion schemes rest on a spectral argument: flow-matching noising progressively masks high spatial frequencies, so above some noise level a down-scaled latent carries *almost* the full surviving signal. We ask whether this licenses a training-side substitution — can the gradient of a low-rank adapter be computed from a downscaled latent at high noise without changing what is learned? We give a first-order account of what the substitution can cost: the gradient gap decomposes into a noise-dependent *input term*, read through the gradient norm, plus a σ -independent *floor* carried by the network’s input–output Jacobian — and we embed the spectral argument, in its strongest tolerance-parameterized form (Xiao et al., 2026), as the null model of the input term: a one-parameter family of per-route boundaries that depends on the downscale *ratio* alone, is compressed to a ~ 0.13 spread in σ on our measured latent spectrum, and admits no floor. Testing this takes metrology, not just measurement: naive finite-draw gap estimators *manufacture* exactly the token-count-ordered floors at issue, so we build a debiased estimator (per-arm self-floors with an attenuation correction, plus draw-count extrapolation) and state every claim at the instrument’s measured resolution. The debiased measurements on a DiT-based flow-matching model reject the null where the accounts diverge: the mildest route’s crossover sits at $\sigma \approx 0.5$, $\sim 3.5\times$ the equal-power prediction; a $2\times$ downscale is certifiably unsafe at *every* noise level, retaining a debiased gap of ≈ 0.30 on pure-noise input; and the surviving endpoint gap equals its image-free control at every route — a floor carried by the compute graph, governed by the *absolute token count* of the downscaled grid, concentrated in early blocks, and separable into a rotary-position-embedding share and a graph residue. Debiasing also cuts both ways: about half of the intermediate route’s raw floor was estimator bias; what remains ($+0.056$, CI $[0.043, 0.071]$) is real but leaves no certifiable safe window at our resolution. The practical output is a measured (route, σ) safety map with explicit detectability thresholds, and an opt-in σ -conditional trainer that substitutes the downscaled latent only on measured-safe steps, for a projected wall-clock ceiling of $\sim 14\%$ at fixed step count on our corpus (end-to-end A/B pending).

1 INTRODUCTION

The cost of a diffusion-transformer forward pass grows superlinearly with resolution: token count grows quadratically with edge length, and attention cost grows again in token count. That makes the title question worth asking precisely. At high noise, can the *gradient* of an adapter be computed on a downscaled latent as a drop-in substitute for the native one — same model, same objective, same σ distribution, only the spatial grid changed on some steps? On our corpus roughly half of all training draws sit at $\sigma > 0.5$, so an affirmative answer on even the mildest downscale is worth a projected $\sim 14\%$ of wall-clock at a fixed step budget (§5).

There is a ready-made answer, and it is the premise of a family of recent methods: under flow-matching noising $z_\sigma = (1 - \sigma)x + \sigma\epsilon$, the per-frequency signal-to-noise ratio falls with σ , and high spatial frequencies drop below the noise floor first — so above a spectral crossover, a downscaled latent carries *almost* the full surviving signal. This observation underwrites scale-wise distillation (Starodubcev et al., 2025), spectral progressive diffusion (Xiao et al., 2026), pyramidal flow matching

(Jin et al., 2024), spectrally-guided noise schedules (Esteves & Makadia, 2026), and progressive-resolution training more broadly (Karras et al., 2018; Chen et al., 2024; Hoogeboom et al., 2023). By the spectral answer, substitution is safe *whenever the noise has masked the frequencies the coarse grid loses* — on our measured latent spectrum, from $\sigma \approx 0.14$.

The measured answer is different in structure, not just in degree. Safety is a per-route boundary jointly governed by two quantities the spectral answer cannot express together: the downscale *ratio* governs a noise-dependent data term, and the *absolute token count* of the target grid governs a σ -independent floor carried by the compute graph itself. The mildest route’s measured boundary sits at $\sigma \approx 0.5 - 3.5\times$ the equal-power prediction, and no tolerance in the spectral family reproduces it — while a $2\times$ downscale is certifiably unsafe at *every* noise level, retaining a gap of ≈ 0.30 on pure-noise input (§4.2, Fig. 1). The evidence collected for the methods above is collected where they operate — at the level of generated output — and a small quality-score difference does not transfer to a small difference in what a training step *learns*.

Answering the question honestly is, unexpectedly, a metrology problem. The natural gap estimator — a redraw floor minus a single demoted-arm cosine — is biased by finite-draw variance, and the bias grows as the demoted grid shrinks: it *manufactures* exactly the token-count-ordered floor signature under test. Every claim in this paper is therefore stated in debiased units, at the instrument’s measured resolution.

The contributions:

1. **Gap metrology: a debiased estimator and a certification resolution** (§3.1, §4.1). We introduce per-arm self-floors with a split-half attenuation correction and a draw-count extrapolation $\text{gap}(D) = \text{gap}_\infty + c/D$ (an uncorrected native self-cosine reads 0.85 at the noise endpoint; corrected, it extrapolates to 1.005, CI [0.994, 1.016]), and derive the instrument’s median certification resolution $\varepsilon^*(N, D)$: “safe” is a one-sided non-inferiority statement at that resolution, never an absolute claim.
2. **A theory of the demotion gap, stated before it is fit** (§3). From the perturbation expansion and cosine geometry we derive, unconditionally, a four-term angular expansion of the gap — data share, graph share, projected interaction, remainder — and, under four named assumptions, its reduction to a two-term account: a route amplitude on a universal mismatch shape read through the U-shaped gradient norm, plus a σ -independent graph floor. The spectral argument, in its strongest tolerance-parameterized form (SPD, Xiao et al., 2026), enters as the null model of the data term’s input-mediated part — ratio-only, and floorless by construction.
3. **Debiased measurements that discharge the theory’s predictions and reject the null where the accounts diverge** (§4). The floor exists, is σ -flat, equals its image-free control at every route, and decomposes into a rotary coordinate-system share and a graph residue; ratio governs the amplitude while absolute size governs the floor; the account *predicts* held-out routes at $\sim 0.07\text{--}0.09$ RMSE, with the floor law hitting two floors it never saw. The exact angular form of the account resolves a systematic overshoot its quadratic reduction leaves (§4.6).
4. **The answer, in deployable form** (§5). A measured (route, σ) safety map at stated resolution — $\{1024 \rightarrow 896 \text{ at } \sigma > 0.5; 1280 \rightarrow 1024 \text{ at } \sigma \gtrsim 0.75\}$, no certifiable window for $1024 \rightarrow 768$, certifiably unsafe everywhere for $\rightarrow 512$ — and an opt-in σ -conditional trainer that substitutes the downsampled latent only on measured-safe steps, for a projected wall-clock ceiling of $\sim 14\%$ at fixed steps (end-to-end A/B pending).

The presentation is theory-first; the practice was observational-first: verdict criteria were frozen before each probe ran, the account was pre-registered with an outcome matrix (including a kill switch) before the discriminating runs, and the debiasing campaign’s decision rules — including the rule under which our own headline floors would have been retracted as variance artifacts — were pre-registered before it ran. Each claim’s epistemic status (pre-registered / confirmatory / post-hoc) is consolidated in the claims ledger (Table 4) rather than hedged sentence-by-sentence.

2 RELATED WORK

Scale-wise and progressive-resolution diffusion. Progressive growing (Karras et al., 2018) and resolution curricula (Chen et al., 2024; Hoogetboom et al., 2023) train at increasing resolution as a schedule. Pyramidal flow matching (Jin et al., 2024) and scale-wise distillation (SwD) (Starodubcev et al., 2025) tie resolution to the noise level itself, with SwD giving the spectral analysis we take as the hypothesis under test: their Fig. 1-style argument shows the noised downsampled latent is close in distribution to a downsampled noised latent above a crossover σ . Spectral Progressive Diffusion (SPD) (Xiao et al., 2026) gives that argument its sharpest available form: a tolerance-parameterized per-frequency activation time under a diagonal Gaussian spectral model — a criterion on the *Bayes-optimal predictor’s* per-band output, not merely on input SNR — from which an inference-time resolution schedule follows; we adopt their Proposition 1 as the null model of our data term (§3.4). SPD also fine-tunes (with LoRA) against *stage-specific modified velocity targets* along the progressive trajectory — changing the objective per resolution stage rather than substituting the latent under the native objective, which is our question; that they observe a training–inference gap motivating this fine-tuning is independently consistent with our finding that resolution changes are not free at the training level. Spectrally-guided noise schedules (Esteves & Makadia, 2026) apply the same per-frequency reasoning to schedule design, deriving bounds on the useful minimum and maximum noise levels from the image spectrum — the spectral account operating as the field’s working model of what noise masks. We are explicit about scope: neither SPD nor SwD *claims* naive gradient equivalence under latent substitution — SwD validates its pipeline at the output level, and SPD fine-tunes precisely because resolution changes are not free. What we test is the tempting extension of their shared premise — that spectral sufficiency of the input licenses a training-side swap — not their methods. Our contribution is orthogonal to all of these: we measure whether input-level closeness transfers to the *training gradient*, holding model and objective fixed — and find it does not transfer at any tolerance.

Flow matching and DiT. We work in the flow-matching / rectified-flow formulation (Lipman et al., 2023; Liu et al., 2023; Esser et al., 2024) on a Diffusion Transformer (Peebles & Xie, 2023) with rotary position embeddings (Su et al., 2024). Adapters are LoRA (Hu et al., 2022).

Positional extension. Position interpolation (Chen et al., 2023), banded/frequency-selective variants (YaRN, Peng et al., 2024), and diffusion-specific dynamic position extrapolation (Issachar et al., 2026; Zhao et al., 2026) rescale rotary coordinates to evaluate a model at unseen grid sizes. We use exact positional interpolation as a *causal intervention* — matching the downsampled grid’s relative phase geometry to native — to decompose the resolution-sensitivity floor, and adopt the σ -gated band schedule of Zhao et al. (2026) (functional form only; their spectral scale laws mispredict our measured boundaries, consistent with our main finding) as a training-time refinement.

Gradient-agreement methodology. Reading training decisions off gradient cosines against a noise-redraw floor follows the general logic of gradient-noise-scale analyses (McCandlish et al., 2018); our estimator operates per σ -bin with means over images, the regime we found reliable (split-half reliability 0.73–0.88), rather than per-image rankings, which were pure noise in our earlier attempts.

3 THEORY: THE DEMOTION GAP AND A TWO-TERM ACCOUNT

This section is definitions and derivations only; every demonstration lives in §4. It defines the estimand and the resolution at which it can be known (§3.1), decomposes the perturbation demotion applies to a training gradient (§3.2), derives the angular expansion of the gap and its conditional reduction to a two-term account (§3.3), states the spectral argument as a null model of one branch (§3.4), and lists the discriminating predictions, each with its designated probe (§3.5).

3.1 ESTIMAND AND RESOLUTION

Objective and noising. With clean latent x , noise $\epsilon \sim \mathcal{N}(0, I)$, and $\sigma \in (0, 1]$, the noised input is $z_\sigma = (1 - \sigma)x + \sigma\epsilon$ and the flow-matching target is $v = \epsilon - x$; the loss is $\|\hat{v}_\theta(z_\sigma, \sigma, c) - v\|^2$ with gradients taken with respect to adapter parameters only.

Demotion. A *demotion* $e_{\text{nat}} \rightarrow e$ replaces the native latent with one produced by pixel-space down-scaling of the source image to the tier- e bucket followed by VAE re-encoding — the ordering SwD validated (“strategy B”); we never downsample in latent space. The caption, the noise draw structure, and the σ under test are held fixed; only the spatial grid changes.

The estimand. For an arm $a \in \{\text{nat}, \text{reenc}, e\}$, let $\bar{g}_a(\sigma) = \mathbb{E}_\epsilon [\nabla_\theta \mathcal{L}]$ be the population per-bin mean adapter gradient of a query (image, caption). The *population demotion distance* is

$$d_e(\sigma) = 1 - \cos(\bar{g}_{\text{nat}}(\sigma), \bar{g}_e(\sigma)) = \frac{1}{2} \|\hat{g}_{\text{nat}} - \hat{g}_e\|^2, \quad (1)$$

with $\hat{g}$ the unit-normalized gradients; $d_e \geq 0$ always. The *reported* quantity of every bin-resolved map and every fit in this paper is the *re-encoding excess* $\Delta_e(\sigma) = d_e(\sigma) - d_{\text{reenc}}(\sigma)$, where the re-encoding control (decode–re-encode at native resolution) prices the VAE round trip that demotion necessarily pays; Δ_e can be negative, d_e cannot. Endpoint and floor tables report debiased estimates of d_e with the control as its own row; each table states which object it reports.

The aggregation operator is part of the estimand. d_e can be aggregated over a probe set per-example (cosine per image, then mean — the batch-1 worst case) or batch-aggregate (gradients summed over the batch before the cosine — the object SGD actually follows). These are different estimands, and no coefficient fitted on one is guaranteed to transfer to the other; the application section publishes one map for each (§5.1).

Resolution: what an instrument can certify. Any estimate of Eq. 1 compares finite-draw estimates of population directions, so every claim carries the instrument’s resolution. For a route whose true excess is zero, the paired estimate $\bar{\Delta}$ fluctuates with standard error $\text{SE}(N, D, \sigma\text{-bin})$, so the one-sided 95% upper bound concentrates near 1.645 SE: no margin smaller than the *median certification resolution*

$$\varepsilon^*(N, D, \text{bin}) = 1.645 \text{SE}(N, D, \text{bin}) \quad (2)$$

is certifiable even for a truly safe route.¹ We *define* “safe at margin ε ” as the non-inferiority statement

$$\text{UB}_{95}(\bar{\Delta}) < \varepsilon, \quad \varepsilon \geq \varepsilon^*. \quad (3)$$

Two consequences hold before any measurement. A “safe” verdict is always a statistical non-inferiority claim at instrument resolution — never a hard bound; a true gap below ε^* is undetectable by construction. And ε^* is the honest currency for negative claims too: “no certifiable window” means the route’s upper bounds never cross below ε at the current (N, D) , which a larger campaign could change — whereas “certifiably unsafe” (lower bound above ε) cannot.

3.2 DATA AND GRAPH PERTURBATIONS

Demotion changes exactly two things about the map from (image, noise, caption) to adapter gradient. Writing the adapter gradient as $g = J^T r$ (input–output Jacobian times prediction residual), the perturbation $\delta g = \bar{g}_e - \bar{g}_{\text{nat}}$ decomposes as

$$\delta g = \underbrace{J^T \Delta \bar{r}(\sigma)}_{\text{data branch } B_e} + \underbrace{\Delta J^T \bar{r}}_{\text{graph branch } C_e} + \underbrace{\Delta J^T \Delta \bar{r} + \dots}_{\text{remainder } R_e}, \quad (4)$$

where:

1. **the data branch B_e :** demotion changes the *data* the objective consumes — spatial-frequency bands above the coarse grid’s Nyquist limit are absent from the input *and* from the target $v = \epsilon - x$ (and the re-encode perturbs what survives) — a perturbation $\Delta \bar{r}(\sigma)$ of the mean prediction residual. Its input-mediated part shrinks as the model’s reliance on input detail shrinks with σ ; its target-mediated part does not: perturbing the image by δx perturbs the residual by $D_z f[(1-\sigma)\delta x] + \delta x$, and the second term carries unit weight at every σ . At $\sigma=1$ the input is pure noise on every grid, so the input-mediated part vanishes in distribution — but the endpoint is *not* data-free by construction; whether the surviving target-mediated share is resolvable is an empirical question, discharged in §4.3.

¹*Median* because at margin exactly 1.645 SE a truly-safe route certifies with only ~50% power; the general margin at size α and power $1-\beta$ is $\varepsilon_{\alpha,\beta}^* = b_D + (z_{1-\alpha} + z_{1-\beta}) \text{SE}$, with b_D a residual finite-draw bias allowance. Eq. 2 is the $\beta = 0.5, b_D=0$ slice.

2. **the graph branch** C_e : token count, rotary phase density, sequence-length-dependent normalization — an operator perturbation ΔJ that is σ -independent, because the grid change does not reference σ .

One qualification is worth stating precisely: J_{nat} and J_e act on different grids, so the two branches are not literal matrix products across arms; they are defined *operationally*, as the two independent perturbation directions in the shared adapter-parameter space where both gradients live, isolated by the designated probes of §4.3.

This decomposition already shows why input-level spectral sufficiency cannot bound the gradient. Input noise masking alone provides no bound on expected-gradient disagreement, because systematic target and graph terms survive noise averaging: the quantity training consumes is $\mathbb{E}_\epsilon[\nabla_\theta \mathcal{L}]$, the target carries the clean image at coefficient one at every σ , and ΔJ never references the spectrum at all. Input-level sufficiency licenses inference-time substitution; it does not extend to training-gradient equivalence.

The graph branch has identifiable parts. Changing the grid changes (i) the rotary coordinate system — every band of the position embedding accumulates phase at a different density across the image — and (ii) non-positional graph statistics: attention softmax over a different token count, sequence-length-dependent normalization, and the sheer capacity of the coarse graph to approximate the fine graph’s computation. Splitting the graph share $\text{Floor}_e = \text{RoPE}_e + \text{Resid}_e$ yields one sharp prediction: the positional share is a *coordinate-system artifact*, so an exact positional-interpolation intervention — evaluating the demoted grid at fractional rotary coordinates that reproduce the native relative phase geometry — should remove RoPE_e wherever the data branch is inactive ($\sigma \rightarrow 1$); and a band-counting sketch says the residue should grow with absolute grid change (the number of rotary bands whose per-extent rotation count crosses decorrelation grows with the extent), attaching an absolute-size governor to Resid_e . We flag the sketch’s epistemic status plainly: like YaRN’s own band thresholds (Peng et al., 2024), its constants are calibrated, not derived; we claim no closed form for Floor_e .

3.3 ANGULAR GEOMETRY: THE FOUR-TERM EXPANSION AND ITS REDUCTION

Cosine geometry fixes how a perturbation is read. With $G(\sigma) = \|\bar{g}_{\text{nat}}\|$ and $P_\perp$ the projector orthogonal to $\bar{g}_{\text{nat}}$, the local identity $1 - \cos(g, g + \delta g) \approx \frac{1}{2} \|P_\perp \delta g\|^2 / G^2$ holds in the small-mismatch limit. Substituting Eq. 4 and expanding the square gives the *four-term angular expansion* — derived unconditionally from the perturbation expansion plus the local cosine identity:

$$d_e(\sigma) \approx \underbrace{\frac{\|P_\perp B_e\|^2}{2G^2}}_{\text{data share } S_e} + \underbrace{\frac{\|P_\perp C_e\|^2}{2G^2}}_{\text{graph share } F_e} + \underbrace{\frac{\langle P_\perp B_e, P_\perp C_e \rangle}{G^2}}_{\text{projected interaction } I_e} + R'_e, \quad (5)$$

where R'_e collects the remainder branch and the beyond-quadratic cosine terms. S_e and F_e are non-negative; I_e can carry either sign — a fact that matters in §4.6.

The reduced two-term account. Under four named assumptions — **(A1)** small remainder: R'_e negligible over the claimed domain; **(A2)** graph-relative stationarity: $\|P_\perp C_e\|/G$ is σ -constant (numerator and denominator share the residual’s scale), so $F_e = \text{Floor}_e$; **(A3)** data-branch factorization: $\|P_\perp B_e\| \approx a_e m(\sigma)$ — a σ -independent route amplitude times a route-independent mismatch shape; **(A4)** negligible projected interaction: $|I_e| \ll S_e + F_e$ over the claimed domain — the expansion reduces to the *two-term account*

$$\text{gap}_e(\sigma) \approx \underbrace{A_e \cdot (m(\sigma)/G(\sigma))^2}_{\text{data term}} + \underbrace{\text{Floor}_e}_{\text{graph term}}, \quad A_e = a_e^2/2, \quad (6)$$

where $m(\sigma) = \|\Delta \bar{r}(\sigma)\|$ is the mismatch of the model’s mean prediction residual across grids (directly measurable, §4.4) and $G(\sigma)$ is the total gradient norm. The quadratic power on m/G is fixed by the cosine geometry *locally*; it is not licensed far from the small-mismatch regime.

The exact angular form. For a perturbation that is purely orthogonal at relative size $\kappa = \|\delta g_\perp\|/G$, the cosine is exactly $1/\sqrt{1+\kappa^2}$, so the corresponding exact gap is

$$d = 1 - \frac{1}{\sqrt{1+\kappa^2}}, \quad \kappa_e(\sigma) = c_e \cdot m(\sigma)/G(\sigma), \quad (7)$$

which saturates instead of overshooting when κ is comparable to one, and whose small- κ limit is Eq. 6’s data term with $A_e = c_e^2/2$. The linear mismatch loading $\|\delta g_\perp\| = c_e m(\sigma)$ is an *empirical* ingredient, not a derivation — so wherever a fitted functional form is reported, Eq. 7 is labeled the *best empirical predictor* and Eq. 6 the *derived small-perturbation form*; §4.6 scores both, held out.

The ledger. Explicitly: the four-term expansion (Eq. 5) is derived unconditionally; the two-term reduction (Eq. 6) is derived *conditional on A1–A4*, each of which is tested by a designated probe (§3.5); the coefficients A_e (or c_e), $m(\sigma)$, $G(\sigma)$, Floor_e are measured, not derived; and any fitted power or loading is labeled empirical. The account was pre-registered with an outcome matrix (including a kill switch: all endpoint gaps ≈ 0 would have falsified the graph term) before the discriminating runs.

3.4 THE SPECTRAL ACCOUNT AS A NULL MODEL OF THE DATA TERM

The spectral argument, made precise, is a model of one part of one branch — and its strongest available form is SPD’s tolerance-parameterized crossover (Xiao et al., 2026). Under the modeling assumption $x_0^{(\omega)} \sim \mathcal{N}(0, P_\omega)$ per frequency ω (with P the clean-latent power spectrum), their Proposition 1 states that for all $\sigma \geq t_\omega$ the *Bayes-optimal* per-band velocity prediction is within tolerance δ of pure noise, $\mathbb{E}[|v^{*(\omega)} - \epsilon^{(\omega)}|^2] \leq \delta$, at

$$t_\omega = \frac{1}{1 + \sqrt{\delta / (P_\omega (1 + P_\omega - \delta))}}, \quad (8)$$

and their Proposition 2 places the safe-substitution boundary for a coarse grid at the t_ω of its Nyquist band. This is a prediction-level criterion, strictly stronger than raw input SNR; the often-quoted equal-power crossover $\sigma_{\text{eq}}(f) = \sqrt{P(f)/(1 + \sqrt{P(f)})}$ is its $\delta = (1 + P_\omega)/2$ slice.

Its scope must be stated exactly: the null constrains only the *input-mediated part of the data branch* B_e . It does not constrain the target-mediated share, the graph branch C_e , the interaction I_e , or the Jacobian gain through which a residual mismatch becomes a gradient mismatch. Read as the null model of the data term, Eq. 8 has three structural properties, each checkable before any gradient is measured:

(N1) One-parameter family. All route boundaries move together under the single tolerance δ ; the family is totally ordered.

(N2) Ratio-only. In normalized frequency the demoted grid’s Nyquist sits at $f_{\text{cut}} = \frac{1}{2} e/e_{\text{nat}}$ — a function of the downscale *ratio* alone. The null therefore cannot distinguish routes of equal ratio at different absolute sizes, and assigns nearly identical boundaries to nearly equal ratios regardless of grid size. It is structurally blind to any absolute-size effect.

(N3) No floor. For every route and every δ , the null’s gap vanishes above a finite boundary: every route becomes safe at some σ .

And under the same diagonal-Gaussian assumption the null extends to the residual level: the posterior mean is a per-band Wiener shrinkage — under a Gaussian prior and Gaussian noise the minimum-mean-square-error estimator is linear with per-band gain $P_\omega/(P_\omega + n_\omega)$, the classical Wiener filter (Wiener, 1949) — so with no cross-frequency coupling the cross-grid mean-residual mismatch is confined to the destroyed band and **(N4)** ordered across routes by destroyed-band energy.

The family is instantiated on our measured latent spectrum in §4.2 (Table 2); its curve-level read in gap units necessarily borrows our bridge — the null emits residual units only, so any gap-unit prediction is *the null transported through the account’s G and a calibrated route gain*, not a parameter-free prediction (§4.6). Two properties of the instantiation are fixed before any gradient is measured: the latents are high-frequency-quiet, so the family is *compressed* — the spread between the mildest and harshest route never exceeds ~ 0.13 in σ at any tolerance — and the two routes our measurements will separate maximally (1024→896, ratio 0.875, and 896→768, ratio 0.857) receive boundaries within 0.01 of each other at *every* δ .

3.5 DISCRIMINATING PREDICTIONS AND FALSIFIERS

The account and the null agree that gaps shrink with σ and order by severity. Everything else is discriminating. Each prediction below names the designated probe that discharges it in §4; several double as falsifiers of the account’s own assumptions.

1. **High- σ persistence** (null: every gap vanishes as $\sigma \rightarrow 1$; account: persistence equal to Floor_e , surviving pure-noise input) — endpoint probe, §4.3. Kill switch: all endpoint gaps ≈ 0 falsifies the graph term.
2. **Endpoint $\equiv$ x-zero $\equiv$ target- α -flat** (the endpoint’s surviving target-mediated share is below resolution — the empirical discharge that the endpoint reads the graph term) — x-zero probe and target-strength sweep, §4.3. This also tests A2’s premise that the floor is a graph property.
3. **Governors** (null: ratio-only, N2; account: ratio governs A_e while absolute target size governs Floor_e) — iso-severity twin and cross-tier transfer, §4.4.
4. **Residual route-uniformity** (null: route-ordered residual mismatch, N4; account’s A3: a route-independent shape $m(\sigma)$) — residual probe, §4.4.
5. **Curve shape** (null: gap shape tied to per-band SNR; account: m/G read — an interior peak where the U-shaped G is smallest, low- G amplification, and a low- σ dip that renormalizing by G removes) — §4.4.
6. **Held-out predictivity** (the reduced account, fit on three routes, must predict routes it never saw from its governors alone) — held-out validation and functional-form refit, §4.6.
7. **The reduction’s domain** (where A2/A4 fail, the reduced form may err in a signed way: a negative I_e can pull the measured gap *below* the additive floor, which no positive two-term form can do) — the designated direct probe is a demote–re-promote arm isolating the data branch with $\text{Floor} = 0$ by construction ([**pending**], §4.6).

4 EVIDENCE

The evidence discharges the theory one term at a time: the instrument and its validation (§4.1), the phenomenon and the null’s boundary-level score (§4.2), the graph floor (§4.3), the data branch and its governors (§4.4), the floor’s decomposition (§4.5), and the account confronted with routes it never saw (§4.6).

4.1 THE INSTRUMENT, VALIDATED

Model and data. All measurements use Anima (CircleStone Labs & Comfy Org, 2025), an open DiT-based flow-matching text-to-image model (2B parameters, 28 transformer blocks, 3D rotary position embeddings, and the Qwen-Image VAE, Wu et al., 2025), fine-tuned with LoRA adapters on an illustration corpus. Images are bucketed by native aspect ratio into resolution *tiers* indexed by nominal edge length $e \in \{512, 768, 896, 1024, 1280\}$; a tier- e image occupies roughly $(e/16)^2$ latent-patch tokens (e.g. $\sim 4,100$ tokens at 1024, $\sim 3,000$ at 896, $\sim 2,160$ at 768, $\sim 1,000$ at 512). The probe adapter is a plain LoRA checkpoint trained at native tiers; we return to this operating-point dependence in §6.

The gradient probe. For each image and each of B uniform σ -bins we accumulate adapter gradients over D stratified noise draws per arm (verdict runs: 8×8 unless noted; $N = 40$ images) and compare arms by cosine similarity of the flattened accumulated gradients:

- **floor**: two independent draw sets at native resolution — the redraw null;
- **reenc**: decode–re-encode at native resolution — the control of §3.1;
- **demote arms**: one per edge e under test;
- **self-floors** (debiased runs): a second independent draw set for *every* arm, giving each arm the same two-estimate structure the native floor has.

The gap is in cosine units (a mismatch *fraction*, scale-free; bin-mean SEM ~ 0.02 at $N=40$). Gap subtraction, not absolute cosines, is the valid read: the floor itself moves with $\|g\|$ across bins. Every bin-mean curve must pass a split-half reliability check over images — a discipline forced on us by an earlier failure in which per-image rankings from the same gradients had split-half reliability indistinguishable from zero. Uniform σ -bins make the training-time σ -density irrelevant to the map; density-weighting the bins reproduces our earlier pooled measurements as a consistency check. Estimator details and the endpoint / x-zero probe modes are in Appendix A.

Finite draws manufacture floors. The naive estimate of Eq. 1 — redraw floor minus a single demoted-arm cosine — is asymmetric: the floor is a cosine between *two* finite-draw estimates of the native gradient, while each demote arm contributes *one* estimate compared against the native mean. Finite-draw noise attenuates every such cosine below its population value, and the attenuation is arm-dependent: per-draw gradient variance grows as the token count falls (the MSE loss averages over fewer elements), so a demoted arm’s single estimate is noisier than the native estimates it is compared to. An iso-direction null — demotion changing *nothing* but the variance — therefore still produces a positive gap that grows as the target grid shrinks: the same signature as “absolute token count sets the floor.” A variance back-of-envelope predicts spurious endpoint gaps of $\approx 0.02/0.05/0.15$ for 896/768/512 — plausibly $\sim 40\%$ of the uncorrected intermediate floors — and a retroactive scan confirms the confound is live: the 1024 $\rightarrow$ 896 endpoint gap decays as $+0.100/+0.035/-0.016$ at $D = 4/8/16$, a clean c/D decay with fitted $c \approx 0.5$. No conclusion in this paper survives on the naive estimator alone.

The native floor itself makes the point sharpest: at $D=8$ –16 the uncorrected endpoint self-cosine reads 0.79–0.85, but a nested draw sweep ($D = 4 \dots 64$) extrapolates it to 1.005 (bootstrap 95% CI [0.994, 1.016]). The population-level per-bin native gradient is a *fixed direction*; the entire redraw floor is finite-draw noise, not intrinsic gradient stochasticity.

The debiased estimator. Two corrections, used together:

1. **Self-floors with attenuation correction.** Every arm runs a second independent draw set g'_d , giving a per-arm self-cosine $\cos_{\text{self}} = \cos(g_d, g'_d)$. The debiased cosine is the classical split-half disattenuation (Spearman, 1904),

$$\hat{c} = \frac{\cos(\bar{g}_{\text{nat}}, \bar{g}_d)}{\sqrt{\cos_{\text{floor}} \cdot \cos_{\text{self},d}}}, \quad \widehat{\text{gap}} = 1 - \hat{c}, \quad (9)$$


which cancels the finite-draw attenuation of numerator and denominator to first order.

2. **Draw-count extrapolation.** Endpoint-mode runs sweep nested draw prefixes $D \in \{4, 8, 16, 32, 64\}$ (the $D=64$ pass contains the smaller draw sets — one set of forwards) and fit $\text{gap}(D) = \text{gap}_{\infty} + c/D$ per route, with a bootstrap CI over images. The two corrections cross-validate: debiased fits are D -flat ($|c| \leq 0.05$ for the 512 route, against $c \approx +0.29$ uncorrected) — the attenuation correction removes exactly the $1/D$ component the sweep isolates.

One caveat delimits where Eq. 9 is readable. At $D=8$ per bin, the *unpaired* debiased estimator is unstable in bins where the floors themselves are small (mid- σ ; the denominator is noisy): there the readable object is the *paired per-image excess* $\Delta_i = \widehat{\text{gap}}_{e,i} - \widehat{\text{gap}}_{\text{reenc},i}$ of §3.1, with non-finite and $|\Delta| > 1.5$ values trimmed — pairing cancels the shared per-image, per-bin draw structure. All verdict maps below report this paired object. Finite-draw cosines are also compiler-kernel-path sensitive, so all pairings live within one run (Appendix A).


Measured resolution. At the verdict grid’s operating point ($N=40$, $D=8/\text{bin}$) the bin-level ε^* (Eq. 2) is 0.03–0.08; endpoint-mode draw sweeps (D up to 64) tighten it to ≈ 0.01 –0.02.

Units and coverage, stated once. From here on, every number in the main text is debiased; the uncorrected historical tables are in Appendix B and are not cited as evidence. Debiasing is not a formality — it moved verdicts in both directions: it *raised* the mid- σ gaps the attenuation had hidden, halved the 768 endpoint floor, and surfaced a small real 896 endpoint gap that single-estimate variance had buried (revision record: Appendix Table 7). Four secondary probes (the 1280-tier iso-severity



(a) Latent RAPSD against the flow-matching noise floor at four noise levels; dashed lines mark the demoted grids’ Nyquist cuts.

(b) The equal-power slice $\sigma_{\text{eq}}(f)$ of the spectral family. Predicted σ^* : 0.136 / 0.146 / ~ 0.20 for demotion to 896 / 768 / 512.



(c) Debiased demotion gap per σ -bin, paired against the re-encoding control (1024-tier natives, $N=40$; values in Table 1). Gray band: bin-level $\pm\epsilon^*$ (Eq. 2); open markers: the $\sigma=1$ endpoint bin. The 896 route enters the band at $\sigma \approx 0.5$; 512 never does.

(d) Estimator context: the redraw-floor cosine and the U-shaped gradient norm $G(\sigma)$ through which every gap is read (§4.4).

Figure 1: The spectral account (top row: computed from the measured RAPSD, predicting safety from $\sigma \approx 0.14$) versus the gradient measurement (bottom row: safe only from $\sigma \approx 0.5$, and only on the mildest route). No tolerance reproduces the measured pattern (Table 2): the equal-power crossover is off by $\sim 3.5\times$ and the $2\times$ downscale never becomes safe.

twin, the 896-tier transfer, the depth split, and the positional-interpolation intervention) predate the debiased instrument; their reads are retained only where one of two defenses applies — the finite-draw bias is strictly positive, so an in-band read is conservative; and paired same-grid contrasts cancel the bias to first order — and their tables are likewise in Appendix B.

4.2 THE PHENOMENON: GAP CURVES AGAINST THE SPECTRAL FAMILY

The spectrum. Anima’s VAE latents are high-frequency-quiet ($P(f) < 1$ above $f \approx 0.16$ cycles/latent-pixel; latent variance 0.42), giving the equal-power crossovers $\sigma^*(896) = 0.136$, $\sigma^*(768) = 0.146$, $\sigma^*(512) \approx 0.20$, with tight per-image spread (p10–p90 ≈ 0.11 – 0.16): the spectral crossover is image-generic, and at that tolerance demotion would be safe over almost the entire σ axis (Fig. 1, top row).

The measurement. Table 1 gives the debiased verdict map for 1024-tier natives demoted to 896/768/512 (the uncorrected historical curves are Appendix Table 6).

Table 1: **Debiased verdict map** (plotted in Fig. 1c): paired per-image debiased excess $\bar{\Delta}$ (SEM) against the re-encoding control, 1024-tier natives ($N=40$, $D=8/\text{bin}$ plus a $\sigma=1$ endpoint bin at the same draws; non-finite and $|\Delta| > 1.5$ per-image values trimmed, retaining 35–40 images per cell). Bin-level ε^* at this (N, D) is 0.03–0.08 (Eq. 2).

σ bin center	0.06	0.19	0.31	0.44	0.56	0.69	0.81	0.94	1.0
896	.091	.155	.208	.032	.013	−.029	.016	.004	.043
(SEM)	.013	.047	.047	.024	.024	.021	.024	.021	.011
768	.149	.238	.324	.206	−.034	−.032	.011	.015	.092
(SEM)	.014	.034	.049	.045	.042	.047	.022	.034	.012
512	.340	.525	.599	.598	.312	.220	.081	.166	.329
(SEM)	.029	.048	.073	.050	.045	.046	.056	.061	.030

Three reads, against criteria frozen before the runs: (i) σ -dependence is real and in the predicted direction, and tier ordering holds at every bin — the *qualitative* spectral intuition survives; (ii) the crossover for the mildest route is $\sigma \approx 0.5$ (the 896 route is clearly gapped through $\sigma \leq 0.31$, marginal at 0.44, and its means sit within ± 0.03 of zero from 0.56 on) — $3.5\times$ the equal-power prediction, and no tolerance rescues the family: the δ that matches this boundary ($\delta = 0.025$, Table 2) predicts 768 safe at 0.52 and 512 safe at 0.62, while SPD’s own published $\delta = 0.01$ predicts all three routes safe by 0.72; (iii) the $2\times$ downscale route is *certifiably unsafe at every noise level*: its paired debiased lower bound stays above ε^* in every bin (means $+0.08$ to $+0.60$), including $+0.17 \pm 0.06$ at $\sigma = 0.94$, where the input is $\sim 97\%$ noise by power in every band the coarse grid cannot represent — the high- σ persistence of prediction 1 in §3.5, which no member of the null family can produce. There is no noise level at which a $2\times$ -downscaled latent yields the native gradient.

Two revisions against the historical record are part of the result. The 768 route’s high- σ plateau was roughly half estimator bias: debiased means sit within noise of zero across $\sigma \in [0.56, 0.94]$, but with upper bounds that never cross below the 0.02 margin at this (N, D) , and a real endpoint floor of $+0.092 \pm 0.012$ remains — so “never safe” softens to *no certifiable window at instrument resolution* (§5.1). Conversely the 896 route’s endpoint shows a small *real* gap ($+0.042 \pm 0.011$) that single-estimate variance had buried; its safe window is the interior $\sigma \in (0.5, 0.94]$, not the endpoint itself.

Ratio does not transfer — the null’s blind spot, measured. Re-running the probe one tier down (896-tier natives, demote to 768/512) fails its pre-registered bar: 896→768 (ratio 0.857, nearly identical to the passing 0.875) retains a high- σ residual roughly $2\times$ the 1024→896 plateau and outside the re-encoding band (Appendix Table 9), and the debiased floors of the *same target grids* reproduce the separation independently — $+0.02$ – 0.04 on the 896-token-count grid versus $+0.06$ – 0.09 on the 768 grid (§4.3). This is exactly property N2’s blindness made concrete: the null assigns these two routes boundaries within 0.01 of each other at every tolerance, and the measurement separates them into *safe* versus *no certifiable window*. “One tier down at high σ ” is not the invariant; §4.4 locates why (the floor is governed by absolute size, and the 768-grid floor alone fills the 0.02 margin).

The tolerance family, scored at the boundary level. Table 2 instantiates Eq. 8 on the measured latent spectrum (P at the 896/768/512 cuts = 0.025/0.029/0.065). Because the latents are high-frequency-quiet, P barely varies across the cuts and the family is compressed: the spread between the mildest and harshest route never exceeds ~ 0.13 in σ at any tolerance. No row reproduces the measured pattern.

Which tolerance? Eq. 8 needs a δ , and it is worth asking where one could legitimately come from. SPD’s published $\delta = 0.01$ is an inference-side choice, selected on a speed–quality Pareto ablation over generated images (Xiao et al., 2026) — nothing ties it to gradients. Machine precision is not the relevant floor either: bf16 rounding on unit-scale per-band amplitudes injects power of order 10^{-5} or below, orders of magnitude beneath any boundary-moving tolerance. Our setting, however, fixes a tolerance *by construction*: the demotion recipe pays a resize–VAE-encode round trip whose latent error has measurable per-band power $D(f)$ — the same pipeline noise the re-encoding control prices at the gradient level — and D is expressed in exactly Eq. 8’s units (per-band power against unit noise). Band content below $D(f_{\text{cut}})$ is invisible to the substitution pipeline even with no resolution

Table 2: The spectral null (Eq. 8) instantiated on the measured latent spectrum: predicted safe boundaries t^* per route as the tolerance δ sweeps its range. SPD’s published default is $\delta = 0.01$; the equal-power slice is the classical crossover. No row reproduces the measured pattern (last line; §5.1): the δ that matches the safe route’s boundary predicts both failing routes safe by $\sigma = 0.63$, and no δ can produce a floor.

tolerance δ	1024→896	1024→768	1024→512	
$(1 + P_\omega)/2$ (equal power)	0.14	0.15	0.20	classical crossover
0.05	0.41	0.43	0.53	
0.025	0.50	0.52	0.62	tuned to the measured 896 gate
0.01	0.61	0.63	0.72	SPD’s published default
0.005	0.69	0.71	0.79	
measured, debiased	≈ 0.5	none certified	unsafe, all σ	floors .02–.04 / .06–.09 / $\approx .3$

Table 3: The floor, debiased. **Endpoint:** at $\sigma=1$ the input is exactly ϵ on every grid, so the input-mediated data term vanishes by construction. **x-zero:** the image is removed from input *and* target; any surviving gap is pure graph-shape sensitivity. Debiased draw-limit values $\widehat{\text{gap}}_\infty$ [bootstrap 95% CI over images] from nested draw sweeps (endpoint: $N=12$, $D = 4 \dots 64$; x-zero: $N=40$, $D = 4 \dots 32$). The endpoint and x-zero columns agree within CI at every route, and the account’s pre-registered kill switch (all endpoint gaps ≈ 0) does not fire.

route	endpoint $\widehat{\text{gap}}_\infty$ [95% CI]	x-zero $\widehat{\text{gap}}_\infty$ [95% CI]
native self-floor (cosine)	1.005 [0.994, 1.016]	—
re-encoding control	−0.003 [−0.017, +0.008]	—
1024→896	+0.019 [+0.010, +0.030]	+0.034 [+0.017, +0.058]
1024→768	+0.056 [+0.043, +0.071]	+0.074 [+0.053, +0.094]
1024→512	+0.304 [+0.197, +0.424]	+0.283 [+0.232, +0.332]

change, so $\delta_{\text{reenc}}(e) := D(f_{\text{cut}}(e))$ anchors the null with zero free parameters. On the probe set this floor and its implied boundaries are **[pending: reenc_noise_floor.py, queued]**; we will report the anchored row alongside the family in Table 2. Either way the confrontation does not hinge on the choice: the structural failures (N2, N3) are δ -independent.

4.3 THE FLOOR IS GRAPH: THREE PROBES, ONE NUMBER

Three designated probes (predictions 1–2 of §3.5) establish the graph term — and with them the premises of the reduction. At the $\sigma = 1$ *endpoint* the input is pure ϵ — identical in distribution across arms — so the input-mediated data term is zero by construction; per §3.2, the target-mediated share need *not* vanish there, so the endpoint alone does not isolate the graph term — the next two probes discharge that step empirically. The debiased endpoint floors are +0.019/+0.056/+0.304 for 896/768/512 (Table 3); the 512 value clears the pre-registered confirmation bar ($\widehat{\text{gap}}_\infty \geq 0.15$) on 12 of 12 probe images individually — the token-count floor is not an estimator artifact.

The *x-zero* probe removes the image from input *and* target everywhere, isolating pure graph-shape sensitivity — and its debiased draw-limit gaps are statistically equal to the endpoint gaps at *every* route (Table 3). The *target-strength sweep* closes the remaining loophole from the α side: re-running the endpoint with target $\epsilon - \alpha x$, $\alpha \in \{0, 0.25, 0.5, 0.75, 1\}$ ($N=12$, $D=16$, one run, shared draws across α), the paired debiased excess is α -flat at the well-conditioned anchors (768: $+0.070 \pm 0.015$ at $\alpha=0$ versus $+0.049 \pm 0.010$ at $\alpha=1$; 512: $+0.269 \pm 0.041$ versus $+0.337 \pm 0.067$; re-encoding control slope +0.003) — no resolvable target-content share. The interior points $\alpha \in [0.5, 0.75]$ are unreadable by construction and excluded: there the native residual $\alpha x - \hat{x}$ passes near cancellation, the gradient norm dips ($60 \rightarrow 33$) and the redraw floor falls to 0.87, so the estimator reads through a small G — the low- G amplification of prediction 5, surfacing along the α axis. Endpoint $\equiv$ x-zero $\equiv$ α -flat: three independent probes converge on one number per route. The endpoint gap *is* the graph floor, as an empirical result — established at our operating point, not assumed by construction.

How does the floor relate to the interior high- σ bins? For 512 the debiased window means (+0.08–0.60) sit at or above the floor, as additivity requires. For 896 and 768 the floors (0.02–0.04 and

0.06–0.09) are *at or below* the interior bins’ certification resolution ($\epsilon^* = 0.03\text{--}0.08$ at $N=40, D=8$): the endpoint sweeps resolve them only because their effective draw budget is an order of magnitude larger. The claim “the high- σ plateau is the floor” therefore sharpens to: the floor is what the plateau converges to when the data term dies and the estimator is given enough draws to see it.

Two further exonerations. The re-encoding control sits at zero within CI — debiased -0.003 $[-0.017, +0.008]$ — so the VAE round trip is not the cause (and, the converse asymmetry: the round trip has *measurable* input-level error, yet its gradient cost is below instrument resolution — input-level error neither implies nor excuses gradient-level cost). And a forward-only probe of the model’s caption-conditioned *prior* across grids shows route distances that are flat where the floors are strongly ordered — so the floor’s route-ordering lives in the Jacobian factor J , not in the model’s resolution-conditioned prior, and not in a training-distribution discontinuity (the never-trained 1280 grid is no more distant from 1024 than heavily-trained adjacent pairs are from each other).

4.4 THE DATA BRANCH: A ROUTE-UNIFORM NUMERATOR OVER A U-SHAPED DENOMINATOR

$m(\sigma)$ is route-uniform — the residual-level null fails. Measuring the mean prediction residual $\bar{r} = \mathbb{E}_\epsilon[\hat{v} - (\epsilon - x)]$ per grid and σ — the exact r in $g = J^T r$ — shows a strong, monotone σ -shape (cross-grid excess $0.89 \rightarrow 0.36$ in relative L_2 from $\sigma=0.125$ to 1) that is the *same for every route* within ± 0.02 , across routes whose gradient floors span 0 to 0.3 and whose crossovers span 0.5 to nonexistent (Appendix Table 10). This kills the diagonal null at the residual level (N4): with no cross-frequency coupling, its mismatch is confined to the destroyed band and ordered by destroyed-band energy, which differs substantially across these routes — route-uniformity says the real mismatch is carried by cross-band structure the diagonal model cannot express. All route identity therefore lives in J , not in the prediction error — a pre-registered null that would equally have falsified the account’s A3 had the residual curves been route-ordered; instead it *establishes* A3’s shape factorization directly. The decay is smooth, Wiener-like posterior shrinkage with no hard gate at the spectral crossover: gap curves read as “collapsed” where the data term sinks below the instrument’s certification resolution, not at σ_{eq} .

$G(\sigma)$ shapes the curve. Because the gap is a cosine it measures a mismatch *fraction*; the total gradient norm is U-shaped in σ ($2.6 \rightarrow 0.4$ at $\sigma \approx 0.3 \rightarrow 9.3$; Fig. 1d) — large at high σ where the model is pulling composition out of the prior and the residual is image-scale, small in the mid- σ refinement regime, and inflated again at low σ by irreducible ϵ . A monotone absolute mismatch divided by a U-shaped norm yields exactly the measured interior-peaked gap curves: bin-mean gap tracks $1/\|g\|$ (Spearman $+0.55$ to $+0.90$ across arms and runs, and within images), and renormalizing by G flips the anomalous low- σ dip into the monotone maximum in every arm of both runs — prediction 5, discharged. This is a post-hoc consistency analysis (so labeled in the ledger): it explains the *shape*; all safety criteria stay in cosine units, since direction is what a normalized optimizer consumes.

The governors: ratio sets the amplitude, absolute size sets the floor. With the floor isolated, what orders the routes? Two probes give a clean division of labor (prediction 3). *Absolute size, not ratio, sets the floor*: the route 1280→1024 is *more* aggressive by ratio (0.80) than 896→768 (0.857, no certifiable window), yet its floor is ≈ 0 and it enters the re-encoding band at $\sigma \gtrsim 0.75$ — refuting any pure ratio rule for the floor, and with it the null’s only degree of freedom. Conversely the floor grows monotonically as the *target* grid shrinks ($0.02\text{--}0.04$ / $0.06\text{--}0.09$ / ≈ 0.30 at target $\sim 3,000/2,160/1,000$ tokens), consistent with “floor = how well the coarse graph approximates the fine graph’s computation.” *Ratio, not absolute size, sets the amplitude*: a synthetic iso-severity route 1280→1120 — edge ratio 0.875 exactly matched to 1024→896, but $1.6\times$ its absolute target size — reproduces the 1024→896 curve bin-for-bin and its crossover $\sigma^* \approx 0.5$, while the same-native 1280→1024 (ratio 0.80) separates and floors later (Appendix Table 8). Amplitude follows ratio; with $\text{Floor}_{1120} \approx 0$ this completes the division: *ratio sets A_e* ; *absolute target capacity sets Floor_e* . The pre-registered contrast (ratio-governed $\Rightarrow \sigma^* \approx 0.5$; capacity-governed $\Rightarrow$ earlier floor) discriminated cleanly. Note what survives of the spectral account here: *ratio is the right governor for the data term* — the null’s error is not its governor but its scope, mistaking one term for the whole gap.

4.5 DECOMPOSING THE FLOOR: DEPTH, AND A POSITIONAL SHARE

Depth, not module type. Splitting the probe gradients per module type (15 groups, including separate rows of the fused QKV projections) and per block (28) shows the floor concentrated in *depth*: early blocks (0–9, peaking at 3–8) carry $\sim 3\times$ the late-block gap, uniformly across every module type within a block (at $512/\sigma=1$, every type sits in 0.22–0.31; Appendix Table 11). The mechanism picture: the first ~ 10 blocks build grid-calibrated token statistics; the divergence propagates into every parameter type’s gradient in those blocks roughly equally, and deeper blocks inherit a washed-out version. Notably, query and key projections show *zero* excess over value projections — which we initially read as refuting a rotary-embedding mechanism. That inference was wrong, in an instructive way.

$\text{Floor}_e = \text{RoPE}_e + \text{Resid}_e$. Landing-side uniformity cannot localize *origin*: a perturbation originating in the position embedding propagates through the block and lands on all module types. The discriminating experiment is the origin-side intervention of §3.2: evaluate the demoted grid with positionally-interpolated fractional rotary coordinates (patch i at $i \cdot H_{\text{nat}}/H_{\text{dem}}$ per axis), making its relative phase geometry exactly match native. At the $\sigma=1$ endpoint this *erases* the 768 floor (paired $\Delta = +0.081$, 78% of images improved) and removes $\sim 30\%$ of the 512 floor (paired $\Delta = +0.096$), while leaving the in-band 896 control untouched (Appendix Table 12; the paired Δ s are same-grid contrasts at matched draws, in which the finite-draw bias cancels to first order). Re-anchored against the debiased floors: the $+0.081$ erasure accounts for essentially all of the 768 floor ($\text{Resid}_{768} \approx 0$); against the 512 floor of ≈ 0.30 , the $+0.096$ erasure leaves a $\text{Resid}_{512} \approx 0.2$ bulk. The mild-route floor is mostly a *coordinate-system artifact*; the harsh-route bulk is a genuine non-positional graph residue (attention-over- N statistics, sequence-length-dependent normalization, capacity). The capacity governor of §4.4 attaches to Resid_e , as the band-counting sketch anticipated. The 896 route’s small floor (0.02–0.04) sits below the PI probe’s paired resolution and remains undecomposed.

The floor as a ledger. Combining the designated probes, the floor decomposes additively, each share measured by its own intervention:

$$\text{Floor}_e = \underbrace{\text{reenc}}_{\approx 0 \text{ by control}} + \underbrace{\text{target content}}_{\approx 0 \text{ by } x\text{-zero} / \alpha\text{-flat}} + \underbrace{\text{RoPE}_e}_{\text{erased by PI at } \sigma=1} + \underbrace{\text{Resid}_e}_{\text{remainder}}$$

with totals 0.02–0.04 (896; undecomposed), ≈ 0.07 (768; RoPE the large majority, $\text{Resid} \approx 0$), and ≈ 0.30 (512; RoPE ≈ 0.10 , $\text{Resid} \approx 0.20$). One caveat on the first share: the re-encoding control (native decode→re-encode) is a *proxy* for the pipeline cost demotion actually pays (source downscale→encode); the demote→re-promote arm of prediction 7 closes that gap empirically and is [pending].

The positional share is a noise-regime artifact only. A σ -resolved follow-up closes the tempting route this opens (training 1024→768 through interpolated coordinates): with image content in the input, the stretched forward is *off-manifold* — the interpolated arm is *worse* than the plain demoted arm through $\sigma \in [0.56, 0.81]$ (paired -0.05 to -0.11) and wins only at $\sigma \geq 0.94$. The floor’s positional share is erasable only where the input is noise-dominated, so the decomposition is a mechanism finding, not a training lever — and the data term of the 768 route is fatal on its own ($+0.09$ – 0.22 over control across the window). A frequency-selective banded alignment does survive in-window (§5.2).

4.6 THE ACCOUNT CONFRONTED: HELD-OUT ROUTES AND THE REDUCTION’S DOMAIN

The predictions so far test the account’s *structure*. Prediction 6 tests whether it is a predictive theory: fit the reduced account on three routes, then predict two routes it never saw, from governors alone. Fit routes: {1024→896, 1024→512, 1280→1120}; held-out: {1024→768, 1280→1024}. Per fit route the excess curve $\hat{\Delta}_e(\sigma)$ is fit with a route amplitude and floor on the measured route-uniform $m(\sigma)$ and each run’s own $G(\sigma)$; two governor models — $A(\text{ratio})$ interpolating the fitted amplitudes, and an exponential floor law $F(n) = F_0 e^{-n/\tau}$ in target tokens n — then generate the held-out predictions with zero per-route freedom. Pass criteria were pre-registered in the analysis script before the numbers.

E5 refit: three functional forms for the data-branch term




Figure 2: The account confronted with held-out routes, under three functional forms for the data term (shared floor structure): the fitted-power form, the derived small-perturbation quadratic (Eq. 6), and the exact angular link (Eq. 7). Left: the severe 512 fit route. Middle/right: the held-out routes 1024→768 and 1280→1024 — every curve is predicted from governors fitted on *other* routes. The exact angular link’s saturation removes the quadratic’s peak overshoot on 1280→1024; the 768 mid- σ dip below the predicted floor is form-independent (see text).

Held-out results. All four pre-registered gates fire: (i) held-out 1024→768: RMSE 0.093 — *better than an oracle quadratic fit on the held-out data itself* (0.105), and beating the null on its committed region ($\sigma > \sigma_{\text{eq}}$, where the null predicts gap 0: RMSE 0.164 null vs 0.096 ours); (ii) held-out 1280→1024: RMSE 0.092 vs oracle 0.056 — inside the pre-registered $2\times$ gate; (iii) the ratio governor holds at the amplitude level: the two ratio-0.875 routes agree at $z = 0.14$ despite a $1.6\times$ difference in target capacity; (iv) the floor law $F(n) = 0.70 e^{-n/1041 \text{ tok}}$, fit on the 512 and 896 floors, predicts $F(4825) = +0.007$ for the 1120 grid (fitted $+0.002$) and $F(2160) = +0.088$ for the held-out 768 grid — against a measured endpoint floor of $+0.092 \pm 0.012$, dead-on. A post-hoc leave-896-out addendum (so labeled) re-derives the governors from 512+1120 alone and predicts the 896 floor at $+0.019$ against the measured $+0.019 [+0.010, +0.030]$: the floor law is 2-for-2 on floors it never saw.

Functional form: the exact link headlines. Refitting the same pipeline under three forms for the data term — a fitted power $A m/G^p$ with p shared and free, the derived quadratic $A (m/G)^2$, and the exact angular link (Eq. 7) — resolves the form question in the geometry’s favor twice over (Fig. 2). First, the free exponent’s own optimum lands on $p = 2.00$ exactly: the “empirical” power form recovers the derived quadratic on its own. Second, the exact link is the best held-out predictor (mean RMSE 0.071 vs 0.093–0.094 for the other two), and its entire margin comes from the 1280→1024 route (RMSE 0.049 vs ≈ 0.095): the quadratic systematically *overshoots* that route’s mid- σ peak, and the exact link’s saturation removes the overshoot — what a first analysis had booked as convexity in the $A(\text{ratio})$ governor was quadratic-form error. The governors are form-invariant (floor law $F(2160) = +0.087\text{--}0.088$ under all three; ratio-twin $z \leq 0.44$). Per the ledger of §3.3: Eq. 7 is the best empirical predictor (empirical through its linear loading $c_e m$), and the derived quadratic is its small- κ limit — one geometry at two orders, not two competing models.

Stated limits. The prediction is *not* within certification resolution anywhere (χ^2/bin 7.6–9.2 held-out under the quadratic): the licensed claim is shape, magnitude class, and governors at $\sim 0.07\text{--}0.09$ RMSE — never “predicts within ε^* .” And the $A(\text{ratio})$ interpolation rests on two distinct ratio values.

The reduction’s domain: one signed failure, and its signature. The 768 route’s mid- σ window is the reduction’s one structural failure — and its shape is diagnostic. Measured excess sits at zero across $\sigma \in [0.56, 0.94]$, *below* the predicted $\text{Floor}_{768} \approx 0.09$: a positive two-term form cannot produce a gap under its own floor, under any coefficients. In the four-term expansion (Eq. 5) this is exactly the signature of a negative projected interaction $I_{768}(\sigma) < 0$ in that window — equivalently, graph-relative stationarity (A2) fails there: the reduced account fails precisely where its named assumptions say it can, and the derived four-term form represents the failure rather than excusing it. The window therefore *delimits the reduction’s domain*; the demote-re-promote probe of prediction 7 — which isolates the data branch with $\text{Floor} = 0$ by construction and reads I_e per bin as $\Delta_{\text{demote}} - \Delta_{\text{repromote}} - \text{Floor}_e$ — is the designated direct test, with the pre-registrable sign $I_{768} < 0$ in-window [pending].

Table 4: Claims ledger: every load-bearing claim, its epistemic status, its designated probe, and where it is discharged. *Pre-registered* = criteria frozen before the run; *confirmatory* = designated probe run under frozen instrumentation, criteria stated with the probe; *post-hoc* = analysis after seeing the data, so labeled.

claim	status	probe	where
Token-count floor exists, carried by the graph	pre-registered kill switch)	(incl. endpoint sweep	§4.3
Endpoint $\equiv$ x-zero $\equiv$ α -flat (no target share)	confirmatory	x-zero, α -sweep	§4.3
896 crossover $\sigma \approx 0.5$; 512 unsafe everywhere	frozen criteria	verdict grid	§4.2
768: no certifiable window (softened from “never safe”)	confirmatory (debi- ased)	verdict grid	§4.2
Ratio governs A_e ; absolute size governs Floor_e	pre-registered contrast	iso-severity twin	§4.4
Residual mismatch route-uniform (kills N4, establishes A3)	pre-registered	residual probe	§4.4
Curve shape = m/G read	post-hoc consistency	renormalization	§4.4
Floor = RoPE + Resid split	confirmatory	PI intervention	§4.5
Account predicts held-out routes (~ 0.07 – 0.09 RMSE)	pre-registered gates	held-out validation	§4.6
Floor law 2-for-2 on unseen floors	gate + post-hoc adden- dum	leave-one-out	§4.6
Exact angular link is the best form	post-hoc (refit)	three-form refit	§4.6
768 mid- σ dip = negative interaction I_e	interpretation; direct	demote-re- promote	§4.6
$\sim 14\%$ wall-clock ceiling	probe pending projected until A/B	fixed-step A/B	§5.2

The null, scored. Against the discriminating predictions of §3.5, the null loses each one it is party to: the boundary pattern is $\{\approx 0.5, \text{no certifiable window, certifiably unsafe everywhere}\}$ (§4.2); gaps persist at pure-noise input, equal to the floor and to the image-free control (§4.3); ratio governs the amplitude while absolute size governs the floor (§4.4); the residual curves are route-uniform under route-ordered gaps (§4.4); and the curve shape follows the m/G read, not per-band SNR. The account also retrodicts the observations that motivated no probe: (i) $\sigma^* \approx 0.5$ rather than 0.14 — the crossover is where the data term sinks below certification resolution, a property of the ratio, so the $3.5\times$ discrepancy is expected, not anomalous; (ii) the 896→768 failure (the 768-grid floor alone fills the certifiable margin — no σ -gate can rescue it); (iii) 1280→1024 flooring *later* ($\sigma^* \approx 0.75$) but cleanly (harsher ratio $\Rightarrow$ larger A_e ; large target $\Rightarrow$ Floor ≈ 0) — the single observation that breaks both pure-ratio and pure-capacity route-ordering stories. A curve-level overlay — each published tolerance δ converted to a predicted gap curve through the account’s bridge (destroyed-band Bayes-residual mismatch, measured G , route gain calibrated on the one safe route) — is **[pending: analysis over existing rows; the anchored line additionally waits on `reenc_noise_floor.py`]**; unit honesty requires saying it will be *the null read through our bridge*, since the null emits residual units only.

A corollary worth stating plainly: the safety of 1024→896 is an *empirical smoothness statement about this network’s function across nearby token counts*, not an information-theoretic statement about the data. There was never a reason to expect a universal ratio invariant.

The claims ledger. Table 4 consolidates every claim’s epistemic status — the paper’s hedging, in one place.

5 THE ANSWER, IN DEPLOYABLE FORM

5.1 THE (ROUTE, σ) SAFETY MAP

Table 5 is the answer to the title question at our operating point: what replaces the spectral rule is not a corrected formula but a measured map at a stated resolution. Safety is a per-route boundary $\sigma^*(\text{route})$, jointly determined by ratio (through A_e) and absolute target size (through Floor_e); neither

Table 5: The measured safety map — *per-example* object (the worst case, what a batch-1 user sees). Floor_e = debiased endpoint floor where the debiased instrument ran; † marks routes measured only by the pre-debiasing instrument (the finite-draw bias is strictly positive, so their ≈ 0 floors and in-band verdicts are conservative reads; §4.1). σ^* = noise level above which the paired debiased excess is non-inferior at instrument resolution (Eq. 3); token ratio = per-forward cost of the demoted grid.

route	edge ratio	target tokens	Floor_e	σ^*	verdict
1024→896	0.875	~3,000	0.02–0.04	≈ 0.5	safe , $\sigma \in (0.5, 0.94]$
1280→1024 †	0.80	~4,100	≈ 0	$\in (0.625, 0.875)$	safe , $\sigma \gtrsim 0.75$
1280→1120 †	0.875	~4,800	≈ 0	≈ 0.5	safe (synthetic grid)
896→768 †	0.857	~2,160	0.06–0.12	—	no certifiable window
1024→768	0.75	~2,160	0.06–0.09	—	no certifiable window
any →512	≤ 0.57	~1,000	≈ 0.30	—	certifiably unsafe, all σ

alone predicts the ordering. For the corpus safe route the fine print is: every bin in $(0.5, 0.94]$ has a debiased mean within ± 0.03 of zero, formal certification at the strict 0.02 margin is achieved at $\sigma = 0.688$ (elsewhere bin-level ε^* exceeds it at the current N, D), and the $\sigma=1$ endpoint itself carries a small real gap ($+0.042 \pm 0.011$) — the safe window is the interior.

The map depends on the aggregation object. Per §3.1, the aggregation operator is part of the estimand, and the two objects genuinely disagree at high σ : pooled-arm probes (pool size 4) collapse the 1024→896 excess to ≈ 0 at every bin $\sigma \geq 0.625$ and the 1024→768 excess to ≈ 0 at $\sigma \geq 0.875$ — the batch-aggregate object SGD follows is more forgiving than the per-example worst case, and the small per-image residual behaves as draw noise that cancels across images rather than as a shared bias. We publish the per-example map as the verdict (it bounds every batch size from below; pooled arms do not yet have self-floors), and a debiased verdict grid at the shipped trainer’s actual batch $\times$ accumulation operating point is designed [**pending**]. A redundancy-stratified analysis finds no per-image demotion-targeting lever (null trend across strata), consistent with the floor being an image-generic network property.

5.2 σ -CONDITIONAL TRAINING ON THE MEASURED-SAFE ROUTE

We wire the one measured-safe corpus route into the LoRA trainer as an opt-in path, exactly as measured — no extrapolation beyond the map:

σ -first draw. The trainer draws each batch’s σ *before* fetching latents, from the unchanged training density (the σ marginal is untouched — demotion conditions on σ , never the reverse). When every sample in the batch is above the gate ($\sigma > 0.5$), the native 1024-tier latent is swapped for a cached demoted-grid sibling produced by the probe’s own measured-safe recipe (pixel-space downscale, VAE re-encode); everything downstream — noise, target, masking, loss — derives from the swapped latent. Off-route images always train native, as does validation. The small real gap at the exact $\sigma=1$ endpoint (§4.3) does not bite here: it is a boundary point of measure zero for the continuous training σ -density, and the last interior bin ($\sigma \in (0.875, 1)$) reads ≈ 0 debiased.

Banded rotary alignment, σ -gated. Motivated by §4.5, a frequency-selective alignment (YaRN-style bands: full positional-interpolation stretch for rotary bands with few rotations across the demoted extent, native spacing for high-frequency bands, ramp between, with the band thresholds scheduled by a sigmoid gate in σ following Zhao et al., 2026) is applied on demoted forwards only. Unlike uniform interpolation it is *not* off-manifold in-window: it passed both pre-registered legs — it erases the static variant’s low- σ liability ($+0.064 \rightarrow +0.033 \pm 0.025$) while preserving its high- σ gains (best and most stable arm at $\sigma \geq 0.59$ and the endpoint, paired -0.05 vs plain demotion). It is a refinement candidate, not a gate-widener: the gate stays $\sigma > 0.5$.

Cost accounting. With $\sim 50\%$ of draws above the gate and the demoted forward at $\sim 0.72\times$ tokens, the fixed-step epoch cost is $\approx 0.5 \cdot 0.72 + 0.5 = 0.86$ — a *projected ceiling* of $\sim 14\%$ wall-clock on our corpus (96% of records are on the safe route), scaling with high-tier content; the realized saving (including cache-load overheads) is measured only by the pending A/B below. The pitch is wall-clock at fixed steps, never “more steps in the same time.”

Weight-space sanity. Under bit-exact deterministic training (paired runs; without determinism, twin runs put a chaos floor of 0.41 on checkpoint-displacement cosine, which would swamp the read), the rotary refinement adds no displacement over the demotion arm it refines (base $\leftrightarrow$ refined 0.319 vs base $\leftrightarrow$ demoted 0.320), with its footprint concentrated in the low-signal early/mid blocks — consistent with §4.5’s depth profile.

What this section does not claim. The end-to-end acceptance gate — a fixed-step A/B at training scale scoring generation-quality non-inferiority (CMMD; Jayasumana et al., 2024) against the realized wall-clock saving — is designed and pre-committed (a regression closes the line) but not yet run; all claims above are gradient-level and weight-space-level. We flag this deliberately: the gap between gradient-level equivalence at $N=40$ and thousands of optimizer steps is exactly the kind of gap this paper is about.

6 LIMITATIONS

One model family, one operating point. All gradient probes ran on one DiT (Anima) with one adapter checkpoint trained at native tiers. A mixed-resolution-trained adapter might equalize (or widen) its own gradients; probing such a checkpoint first is the designated bound on how far the map generalizes. The floor’s dependence on the fine-tuning operating point through J is untested (the designated probe: fine-tune at the 1280 tier, re-probe 1280 $\rightarrow$ 1024). Relatedly, the probe pool’s relationship to the adapter’s own fine-tune set was never a controlled axis (2 of 40 verdict stems are in it); a controlled-adapter 2×2 factorial — two LoRAs trained on opposite style clusters under frozen stem-level membership manifests, probed on both clusters — is designed to measure this dependence rather than concede it, and its adapter training is in progress at the time of writing [pending].

Gradient-level, not end-task. The map is built from accumulated per-step gradients; integration over an optimizer trajectory is addressed only by the pending A/B above.

Account resolution and domain. The account predicts held-out routes at $\sim 0.07\text{--}0.09$ RMSE — never within ε^* (§4.6); its $A(\text{ratio})$ governor rests on two distinct ratio values; the reduction’s domain excludes the 768 mid- σ window pending the direct interaction probe; and the band-counting sketch for $\text{RoPE}_e/\text{Resid}_e$ is calibrated, not derived.

Instrument resolution. “Safe” is one-sided non-inferiority at the measured certification resolution (§3.1): $\varepsilon^* = 0.03\text{--}0.08$ per bin at the verdict grid’s (N, D) , $\approx 0.01\text{--}0.02$ at the endpoint sweeps. True gaps below ε^* are undetectable by construction, the σ^* boundaries are localized only to a bin, and every “no certifiable window” claim is (N, D) -relative — a larger campaign could certify (or refute) the 768 high- σ window whose debiased means already sit near zero.

Debiased coverage. The debiasing campaign covered the corpus verdict grid: the endpoint draw sweeps, the 8-bin map, x-zero, and the α -sweep. The 1280-tier iso-severity, 896-tier transfer, depth-split, and PI-intervention probes remain pre-debiasing reads, defended where used by the one-signed-bias argument and by paired same-grid contrasts (§4.1); pooled arms have no self-floors yet. The 1280-tier curves enter the held-out validation with this caveat recorded.

Scope. Adapter (LoRA) gradients on a frozen backbone; full fine-tuning may differ. Pixel-space demotion only (latent-space downsampling was not tested here; SwD found it inferior). The σ -gated banded-alignment parameters were taken from one shot at pre-registered values, not searched.

7 CONCLUSION

When does training on downsampled images yield the same gradients? Not when the noise has masked the frequencies the coarse grid loses — the spectral answer, even in its strongest tolerance-parameterized form, is a model of only one part of one branch of what demotion perturbs, and the measurements reject it at every point of structural divergence: the mildest route’s crossover sits at $3.5\times$ the equal-power prediction, no tolerance reproduces the map, and a $2\times$ -downsampled latent is certifiably unsafe at every noise level, because the network function itself is grid-calibrated — partly through its positional coordinate system, partly through an irreducible dependence of the computation on token count. The measured answer has two conditions, matching the two terms of the account: the target grid must be large enough in *absolute token count* that its graph floor sits below certifica-

tion resolution, and the noise level must exceed a per-route boundary where the ratio-governed data term dies. On our model that means $1024 \rightarrow 896$ at $\sigma > 0.5$ and $1280 \rightarrow 1024$ at $\sigma \gtrsim 0.75$ — and never a $2\times$ downscale, at any noise level. The answer is predictive, not merely descriptive: the exact angular form of the account, with its two measured governors, predicts held-out routes at $\sim 0.07\text{--}0.09$ RMSE, and its floor law hit two floors it never saw. Getting there required treating the question as a metrology problem — naive finite-draw estimation manufactures exactly the token-count-ordered floor signature under test, half the intermediate route’s raw floor was estimator bias, and our own headline claims were re-derived, and in one case softened, under the correction. On the one route our corpus makes safe, the map projects a ceiling of $\sim 14\%$ wall-clock at fixed steps, pending the end-to-end A/B. We offer the account and its instrument — per- σ -bin gradient agreement with a redraw floor, a re-encoding control, per-arm self-floors, and a stated certification resolution, criteria frozen before data — as the cheap, general test that any future “train part of the time at lower resolution” proposal should pass before it is believed; without the debiasing, that test manufactures the very effect it looks for.

REPRODUCIBILITY STATEMENT

The probe scripts (`run_sigma_probe.py` — including the `--self_floor` and `--draw_sweep` debiasing modes — `run_prior_distance.py`, `rapd.py`, `compare_ckpt_dw.py`), the held-out validation and refit analyses (`e5_holdout.py`, `e5_refit.py`), the trainer wiring (`--sigma_lowres`), and the frozen pre-registration documents (including the debiasing campaign’s decision rules) are in the public repository at https://github.com/sorryhyun/anima_lora (project line `project/sigma_lowres/`). The debiased verdict runs’ complete per-image records (`result.json` + `per_image.jsonl`) ship in-repo under `project/sigma_lowres/paper_bench/runs/`; the earlier raw-instrument run archives are excluded by the repository’s results ignore rule and will be published as a companion tarball **[pending]** — until then, raw-run numbers are reproducible from the scripts but their original records are not themselves public. Paired training runs use a `--deterministic` mode (deterministic attention backward included) verified bit-identical across twin runs; finite-draw cosines are kernel-path-sensitive, so all cosine pairings are computed within a single run (Appendix A).

REFERENCES

- Junsong Chen, Jincheng Yu, Chongjian Ge, Lewei Yao, Enze Xie, Yue Wu, Zhongdao Wang, James Kwok, Ping Luo, Huchuan Lu, and Zhenguo Li. PixArt- α : Fast training of diffusion transformer for photorealistic text-to-image synthesis. In *International Conference on Learning Representations*, 2024.
- Shouyuan Chen, Sherman Wong, Liangjian Chen, and Yuandong Tian. Extending context window of large language models via positional interpolation. *arXiv preprint arXiv:2306.15595*, 2023.
- CircleStone Labs and Comfy Org. Anima: an open text-to-image model for non-photorealistic styles. <https://huggingface.co/circlestone-labs/Anima>, 2025.
- Patrick Esser, Sumith Kulal, Andreas Blattmann, Rahim Entezari, Jonas Müller, Harry Saini, Yam Levi, Dominik Lorenz, Axel Sauer, Frederic Boesel, Dustin Podell, Tim Dockhorn, Zion English, Kyle Lacey, Alex Goodwin, Yannik Marek, and Robin Rombach. Scaling rectified flow transformers for high-resolution image synthesis. In *International Conference on Machine Learning*, 2024.
- Carlos Esteves and Ameesh Makadia. Spectrally-guided diffusion noise schedules. *arXiv preprint arXiv:2603.19222*, 2026.
- Emiel Hooeboom, Jonathan Heek, and Tim Salimans. Simple diffusion: End-to-end diffusion for high resolution images. In *International Conference on Machine Learning*, 2023.
- Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022.

- Noam Issachar, Guy Yariv, Sagie Benaim, Yossi Adi, Dani Lischinski, and Raanan Fattal. DyPE: Dynamic position extrapolation for ultra high resolution diffusion. In *International Conference on Machine Learning*, 2026. arXiv:2510.20766.
- Sadeep Jayasumana, Srikumar Ramalingam, Andreas Veit, Daniel Glasner, Ayan Chakrabarti, and Sanjiv Kumar. Rethinking FID: Towards a better evaluation metric for image generation. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024.
- Yang Jin, Zhicheng Sun, Ningyuan Li, Kun Xu, Kun Xu, Hao Jiang, Nan Zhuang, Quzhe Huang, Yang Song, Yadong Mu, and Zhouchen Lin. Pyramidal flow matching for efficient video generative modeling. *arXiv preprint arXiv:2410.05954*, 2024.
- Tero Karras, Timo Aila, Samuli Laine, and Jaakko Lehtinen. Progressive growing of GANs for improved quality, stability, and variation. In *International Conference on Learning Representations*, 2018.
- Yaron Lipman, Ricky T. Q. Chen, Heli Ben-Hamu, Maximilian Nickel, and Matt Le. Flow matching for generative modeling. In *International Conference on Learning Representations*, 2023.
- Xingchao Liu, Chengyue Gong, and Qiang Liu. Flow straight and fast: Learning to generate and transfer data with rectified flow. In *International Conference on Learning Representations*, 2023.
- Sam McCandlish, Jared Kaplan, Dario Amodei, and OpenAI Dota Team. An empirical model of large-batch training. *arXiv preprint arXiv:1812.06162*, 2018.
- William Peebles and Saining Xie. Scalable diffusion models with transformers. In *IEEE/CVF International Conference on Computer Vision*, 2023.
- Bowen Peng, Jeffrey Quesnelle, Honglu Fan, and Enrico Shippole. YaRN: Efficient context window extension of large language models. In *International Conference on Learning Representations*, 2024.
- Charles Spearman. The proof and measurement of association between two things. *The American Journal of Psychology*, 15(1):72–101, 1904.
- Nikita Starodubcev, Ilya Drobyshevskiy, Denis Kuznedelev, Artem Babenko, and Dmitry Baranchuk. Scale-wise distillation of diffusion models. *arXiv preprint arXiv:2503.16397*, 2025.
- Jianlin Su, Murtadha Ahmed, Yu Lu, Shengfeng Pan, Wen Bo, and Yunfeng Liu. RoFormer: Enhanced transformer with rotary position embedding. *Neurocomputing*, 568:127063, 2024.
- Norbert Wiener. *Extrapolation, Interpolation, and Smoothing of Stationary Time Series*. MIT Press, 1949.
- Chenfei Wu, Jiahao Li, Jingren Zhou, Junyang Lin, Kaiyuan Gao, et al. Qwen-Image technical report. *arXiv preprint arXiv:2508.02324*, 2025.
- Howard Xiao, Brian Chao, Lior Yariv, and Gordon Wetzstein. Spectral progressive diffusion for efficient image and video generation. *arXiv preprint arXiv:2605.18736*, 2026.
- Bingxuan Zhao, Qing Zhou, Yu Wang, Chuang Yang, and Qi Wang. σ : Sigmoid modulation for ultra high resolution diffusion. In *International Conference on Machine Learning*, 2026.

A INSTRUMENT DETAILS

Estimator. Per image, arm, and σ -bin, adapter gradients are accumulated over D stratified draws and flattened; cosines are taken between accumulated vectors. Per-bin cosines at $D=8$ are not comparable *across* bins (the floor drops where $\|g\|$ is small); the gap subtraction against the same-bin floor is the valid read. Split-half reliability over images of each bin-mean curve is mandatory (observed 0.73–0.88 on verdict runs); a per-image variant of the same quantity has split-half reliability ≈ 0 and is never used.

Controls. The re-encoding arm prices the VAE decode–encode round trip that demotion necessarily pays; in pre-debiasing runs its gap served as the validity band (± 0.04 ; observed $|\cdot| \leq 0.054$ across all bins of the main runs, per-module-group ≤ 0.045 in the split runs), a role superseded by the paired debiased read of §4.1. The redraw floor prices finite-draw noise. Density-weighting the uniform bins by the trainer’s σ -density reproduces earlier σ -marginalized measurements from the same instrument family (consistency check).

Debiasing (§4.1), implementation. Under `--self_floor`, every arm (re-encoding and each demote variant) runs a second independent draw set from a disjoint seed stream, giving per-arm self-cosines alongside the native floor; Eq. 9 is then computed per image and bin, and the map read is the paired excess $\Delta_i = \widehat{\text{gap}}_{e,i} - \widehat{\text{gap}}_{\text{reenc},i}$ with non-finite values (negative self-cosines under the square root at low D) and $|\Delta_i| > 1.5$ trimmed (0–5 of 40 images per cell). Under `--draw_sweep`, draw prefixes are nested ($D = 64$ contains the $D \leq 32$ sets; prefix sums, no extra forwards) and $\text{gap}(D) = \text{gap}_\infty + c/D$ is fit on route means with a bootstrap CI over images. Finite-draw cosines are kernel-path chaotic: twin runs sharing a warm compiler kernel cache agree to $|\Delta \cos| \leq 0.015$, while a run compiled to a different kernel set lands up to ≈ 0.3 away at $D=2$ — so gap/floor/self-floor pairings are never compared across processes, a guarantee the single-run instrument provides by construction; a `--deterministic` mode (bit-exact across twin runs) covers the paired training A/Bs.

Endpoint and x-zero modes. The $\sigma=1$ endpoint bin feeds input exactly ϵ (the input-mediated data term vanishes by construction). The x-zero mode zeroes x in input *and* target on every grid, keeping captions and exact demoted latent shapes; with the $(1-\sigma)x$ term absent, low- σ bins are off-manifold and the $\sigma=1$ read is primary. In that regime the residual is $\approx -\hat{x}_{\text{prior}}$, so “graph-dominated” includes the model’s grid-conditioned prior; the prior-distance probe (§4.3) subsequently dissociated the prior from the floor’s route ordering.

Held-out validation and refit (§4.6), sources. The fit consumes the debiased paired curves of the 1024-tier verdict run, the 1280-tier curves (pre-debiasing paired reads; caveat recorded in §6), the route-uniform $m(\sigma)$ from the residual probe, and each run’s own bin-mean $G(\sigma)$ (gap and G always from the same process, per the kernel-path rule). The shared-power fit scans $p \in [1, 2]$ jointly across fit routes with per-route weighted least squares; the exact-link fit is a per-route grid with profile-likelihood uncertainties on c_e . Oracle baselines are quadratic-in- σ fits on the held-out data itself (the noise ceiling for a 3-dof curve); the null baseline is $\text{gap} = 0$ on the null’s committed region $\sigma > \sigma_{\text{eq}}$.

Spectral-null computation. Table 2 evaluates Eq. 8 at the demoted grids’ Nyquist frequencies $f_{\text{cut}} = \frac{1}{2} e/e_{\text{nat}}$ on the $N=40$ mean radially-averaged latent power spectrum (64 radial bins; P at the 896/768/512 cuts = 0.025/0.029/0.065). The equal-power slice $\delta = (1 + P_\omega)/2$ recovers $\sigma_{\text{eq}} = \sqrt{P}/(1 + \sqrt{P})$ exactly. Across $\delta \in [0.001, 0.5]$ the mildest–harshes spread never exceeds 0.13, and the 1024→896 versus 896→768 boundaries (cut frequencies 0.4375 vs 0.4286) never separate by more than 0.01. The parameter-free anchor δ_{reenc} of §4.2 is measured by the same estimator: RAPSD of the re-encode arm’s latent error (fresh resize–encode minus cached latent, the gradient probe’s own control chain) on the same probe set, interpolated at each route’s cut frequency [**pending: reenc_noise_floor.py**].

B HISTORICAL (PRE-DEBIASING) TABLES

The tables in this appendix are the raw-estimator record: retained for continuity and for the reads defended in §4.1 (one-signed bias; paired same-grid contrasts), never cited as debiased evidence.

Table 6: **Raw estimator** (historical record; superseded by Table 1): demotion gap by σ -bin, 1024-tier natives ($N=40$, bin-mean SEM ~ 0.02 , re-encoding control $|\text{gap}_{\text{reenc}}| \leq 0.054$ everywhere, split-half reliability 0.73–0.83). Bold = within the re-encoding band. Raw values understate mid- σ gaps (the floor attenuates the read) and overstate endpoint floors (single-estimate variance bias); the debiased map corrects both. The $\cos_{\text{floor}}$ and $\|g\|$ rows are the estimator context plotted in Fig. 1d.

σ bin center	0.06	0.19	0.31	0.44	0.56	0.69	0.81	0.94
gap ₈₉₆ (ratio 0.875)	.110	.162	.148	.137	.048	.030	.030	.053
gap ₇₆₈ (ratio 0.75)	.144	.216	.208	.223	.164	.163	.115	.063
gap ₅₁₂ (ratio 0.5)	.348	.410	.355	.469	.430	.391	.289	.296
$\cos_{\text{floor}}$	.84	.65	.51	.65	.70	.77	.80	.83
$\ g\ $ (native)	2.6	0.5	0.4	0.5	0.7	1.1	2.0	9.3

Table 7: Raw endpoint reads ($D=16$, $N=40$) beside their debiased draw-limit values (Table 3) — the debiasing’s revision record: the 768 floor halves, the 896 floor surfaces.

route	raw endpoint ($D=16$)	endpoint $\widehat{\text{gap}}_{\infty}$ [95% CI]
native self-floor (cosine)	0.85	1.005 [0.994, 1.016]
re-encoding control	−0.038	−0.003 [−0.017, +0.008]
1024→896	−0.009	+0.019 [+0.010, +0.030]
1024→768	+0.127	+0.056 [+0.043, +0.071]
1024→512	+0.326	+0.304 [+0.197, +0.424]

Table 8: Iso-severity probe (**raw estimator**): gap by σ -bin on the 1280-tier cache ($N=24$, 4 draws/bin) against the corpus 1024→896 curve. The ratio-matched routes coincide within ~ 1.4 combined SEM at every bin despite a $1.6\times$ difference in absolute target size; the same-native harsher-ratio route separates. *896’s canonical 16-draw endpoint is −0.009 raw, +0.019 debiased (Table 3).

σ	0.125	0.375	0.625	0.875	1.0
re-encoding control	+0.047	−.005	+0.048	+0.020	−.012
1280→1120 (ratio .875)	.129	.110	.054	−.012	−.007
1280→1024 (ratio .80)	.170	.178	.111	.002	.061
1024→896 (ratio .875)	.132	.076	.077	.049	.100*

Table 9: Ratio-transfer run (**raw estimator**): 896-tier natives ($N=40$), pre-registered bar “gap₇₆₈ within the re-encoding band at $\sigma \geq 0.5$ ” — **fails** (residual 0.06–0.12, $\sim 2\times$ the 1024→896 plateau; $\sim$ half of it is expected estimator bias, §4.1 — the separation is independently confirmed by the debiased same-grid floors, §4.3).

σ bin	0.06	0.19	0.31	0.44	0.56	0.69	0.81	0.94
896→768	.114	.209	.168	.143	.124	.056	.092	.061
896→512	.318	.372	.308	.340	.320	.280	.329	.217

Table 10: Route-uniformity of the mean prediction residual: cross-grid excess of $\bar{r} = \mathbb{E}_{\epsilon} [\hat{v} - (\epsilon - x)]$ over split-half floors (relative L_2 ; same-grid re-encoding control ≤ 0.02 everywhere). The three main routes coincide within $\sim \pm 0.02$ at every σ while their gradient floors span 0 to 0.3.

σ	0.125	0.375	0.625	0.875	1.0
1280→1024	.885	.825	.715	.465	.360
1024→896	.862	.808	.706	.459	.378
896→768	.860	.814	.715	.466	.393
768→512	.897	.872	.767	.508	.435

Table 11: Depth localization of the floor (**raw estimator**): mean per-block gap at $\sigma=1$, early = blocks 0–9, late = 14–27, for the endpoint (ep) and x-zero (xz) probes. Early blocks carry $\sim 3\times$ the late-block gap. The apparent content share (ep – xz) is a late-block minority effect in raw units; the debiased draw-limit comparison finds ep = xz within CI at the whole-adapter level (Table 3).

route	ep early	ep late	xz early	xz late
1024→512	.357	.223	.351	.125
1024→768	.164	.085	.121	.033
1024→896	.023	.015	.044	.012

Table 12: Positional-interpolation intervention at the $\sigma=1$ endpoint ($N=40$, 16 draws; **raw** arm values — the paired Δ column is a same-grid contrast in which the finite-draw bias cancels to first order): exact phase-geometry alignment erases the mild-route floor and $\sim 30\%$ of the harsh-route floor.

route	plain (SEM)	PI-aligned (SEM)	paired Δ (SEM)	improved
1024→896 (control)	−0.021 (.041)	−0.040 (.040)	—	—
1024→768	+0.080 (.048)	−0.001 (.039)	+0.081 (.031)	78%
1024→512	+0.320 (.058)	+0.224 (.056)	+0.096 (.039)	70%